

A short reflection (250–350 words) explaining how altering the x value impacted the output of your $P+x$ version. Additionally, discuss how you would implement a $P+7$ technique in which all nouns are replaced with their seventh-highest probability alternatives.

GPT-2 to begin with is not a very good model for predicting poetry, so the initial trial for x probability substitutes were not impressive. If x rose to some high number, it was also likely to create garbage. However, this is an interesting point to notice. Most average users never see the low probability responses because models don't run on high heat. The main reason for this, from what I know, is reliability which goes against the incentives in the context of poetry. Some may say it's because chaos and creativity go hand in hand, but realistically, an artist does have some predictability to what they do. It's just that most models can justify training for a specific scenario that isn't commercially viable, leaving the real problem of such models obvious: they can't adapt to a new type of prediction easily, and need to default to some kind of average case.

This ties into our next problem of replacing every noun with a low probability alternative. There are multiple ways of doing it. You could create the text as normal, but pass over each noun in the context of what was written and replace the results. You could also actively impose a selector, which picks low probability outcomes at every moment where a noun could be selected. In either case you would reduce the poem to a garbled mess.